

Techniques and tools for measuring energy efficiency of scientific software applications

*16th international workshop on Advanced Computing and Analysis
Techniques in Physics Research*

Giulio Eulisse Fermi National Accelerator Lab. (US)

Tapio Niemi Helsinki Institute of Physics (FI)

Peter Elmer Princeton University (US)

David Abdurachmanov Vilnius University (LT)

Goncalo Pestana Aalto University (FI)

Filip Nybäck Filip Aalto University (FI)

Jukka K. Nurminen Aalto University (FI)

Zhonghong Ou Aalto University (FI)

Robert Knight Princeton University (US)

Gonçalo Pestana

1. Why is energy efficiency important on High Throughput Computing (HTC) ?

2. Measuring power consumption

Where is energy used ?

External measurements

Chip monitors

Software based measurements

3. Study case

External measurements

Chip monitors

IgProf for energy measurements

Results: ARM vs Intel comparison

4. Conclusions

1. Why is energy efficiency important on HTC ?



Worldwide LHC computing grid estimated 10 MW of power for 350k x86 cores

only CMS experiment will increase 2-3 orders of magnitude the amount of data in next 15 years

Future: Scalability, raising costs ?

next: where ?

2.1. Where is the energy used ?

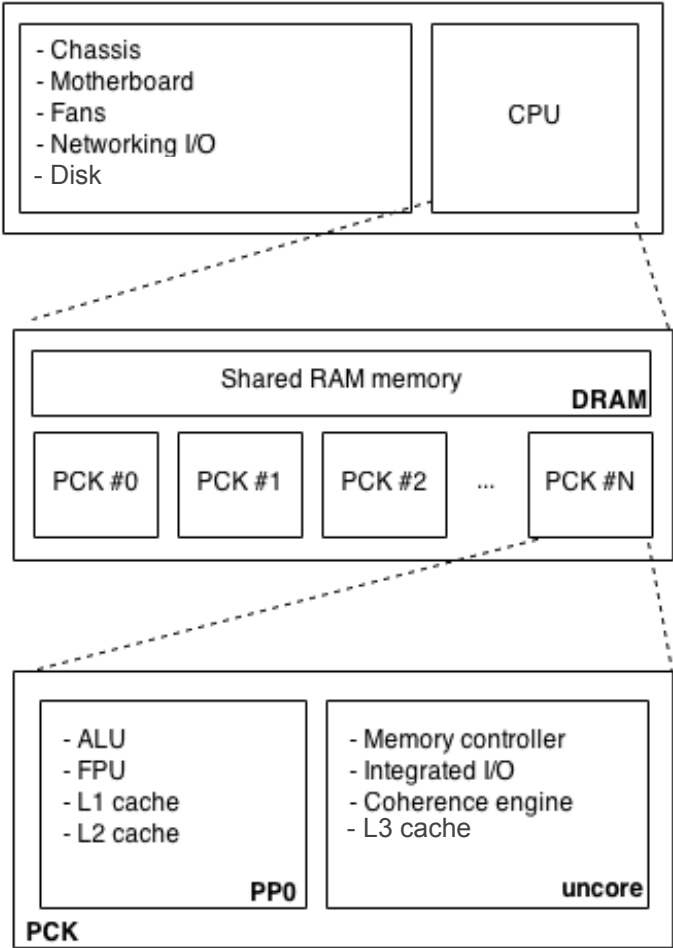
HTC systems are **complex**



Several components

Different **granularities**

CPU can be optimised
more control than other components



2.2. Where is energy used ? > External measurements

Account for the **whole system**

Power consumed without breaking down the system in components

Accuracy: around $\pm 3\%$ for clamps

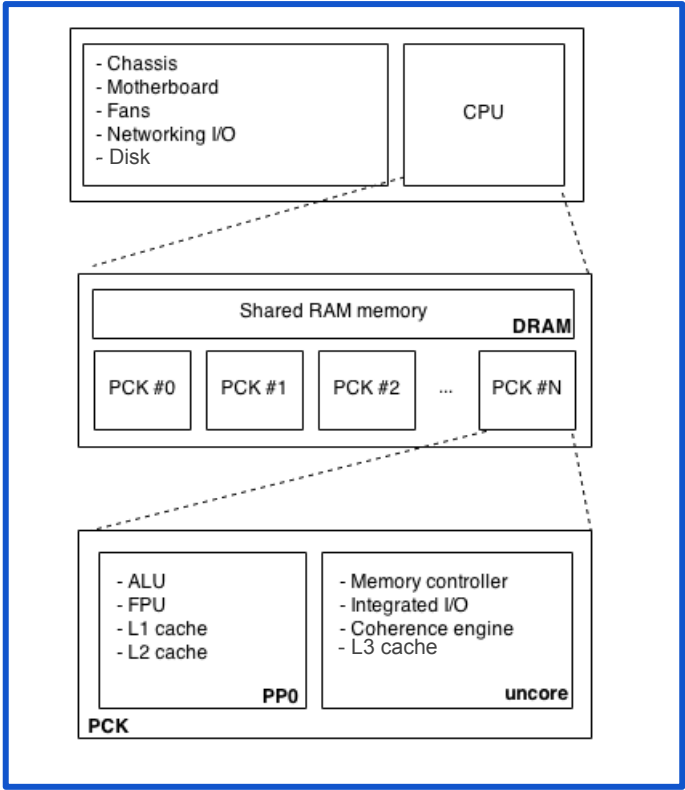
Resolution order of seconds (Hz)



<http://www.p-mastech.com>



<http://www.racksandpower.com>




next: internal

2.2. Where is energy used ? > *Chip monitors*

Monitor chips on the motherboard that measure energy consumed by different components of CPU

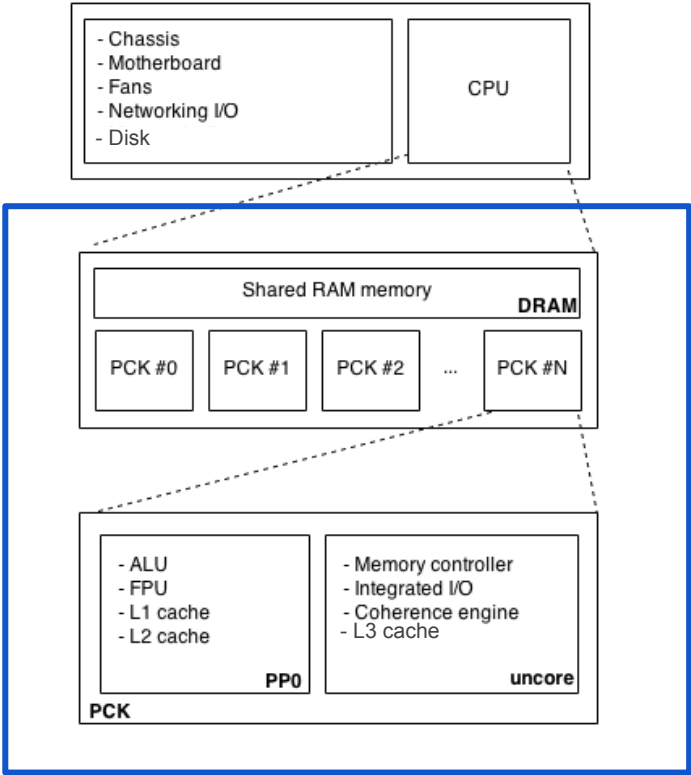
There is no standard solution.

Accuracy: *high* (ex: 0.5% gain error)
Resolution order of microseconds (MHz)



Texas Instrument (TI) INA 231 monitoring chip on ODROID XU+E

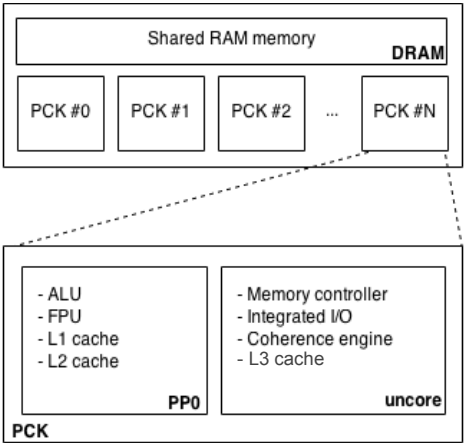
<http://hardkernel.com>



Running Average Power Limiting (RAPL) by Intel

Provides a platform for power monitoring and power limiting of SoC.

Different sampling **domains**
package (*PKG*), DRAM, core



- Low level measurements** package, cores, dram
- Resolution** according to Intel, sampling frequency up to ~1 kHz
- Power capping** is also supported by RAPL
- Accuracy** high (according to *Intel*)

2.2. Where is energy used ? > *Chip monitors* > **RAPL**


Example of RAPL

System with 4 sockets

Sockets #0 and #1 working

Sockets #2 and #3 idle

Possible to understand how
packages, cores & dram consume
energy



2.2. Where is energy used ?

Recap

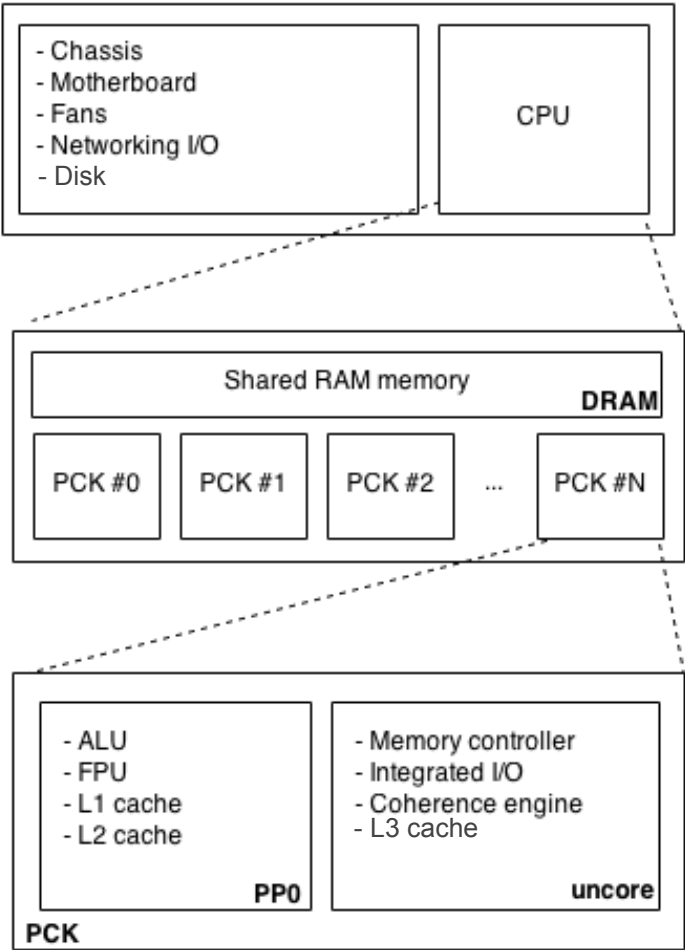
Energy consumption in HTC is an issue

HTC systems are complex

Different tools for different granularities and components:

External measurements

Chip monitors



2.2. Where is energy drained ? > Software-based> IgProf

IgProf

application profiler developed at CERN by the CMS software team

general purpose. open source. not experiment specific

measures performance (time spent in functions) and memory usage at *runtime*

allows developer to understand:

- bottlenecks

- where code needs to be optimised

cross platform: recently ported to 64-bit ARM, also supports 32-bit ARM, Intel x86 and x86-64

2.2. Where is energy drained ? > Software-based> IgProf

IgProf & energy profiling

Uses RAPL and PAPI to measure energy consumed.

Map functions and low level operations with **energy consumption**

more info (strategies, results, examples)

***paper** and <http://igprof.org/>*

3. Study case

Comparison ARMv7 vs Intel XEON

32bit ARMv7 is used on smartphones. comparison with Intel XEON

measurements

Internal

RAPL for *Intel*

cross platform *chip monitor* integrated (TI INA 231) for ARMv7

External

workload

ParFullCMS

Multithreaded

Geant4 benchmark application

Uses the CMS geometry

3. Study case

ARMv7 Exynos5 Octa Cortex™
4x A15 @ 1.6Ghz and/or A7 cores (big.LITTLE technology)
2 GB RAM
ARMv7/32bit
development board

Intel 32x Intel™ Xeon™ CPU E5-2650 @ 2.00GHz
252 GB RAM
system on a rack



3. Study case

What to expect

ARMv7 performs better (energy efficiency perspective)

Both ARM and Intel do not perform better when overcommitted (*more threads than number cores*)

Power consumed ratio external/internal is bigger on Intel (*rack vs dev. board*)

Chip monitors and external measurement techniques in action


Chip monitors measurements

TI INA 231

- core + dram

RAPL

- core + dram




3. Study case > External measurements

External measurements


External power monitor

rack PDU power monitor




3. Study case > External measurements

ARM ratio *CPU+mem* vs overall



Intel ratio *CPU+mem* vs overall



3. Study case > Comparison ARM vs Intel

32bit ARMv7 with interesting results.

Different system's components → different tools for measurement

External, ***chip monitors*** and **software-based** measurements give good picture of the system's energy consumption at different levels.

4. Wrapping up

Energy efficiency is a major concern of HTC

Important to be aware of where and how energy is consumed on a HTC system

There are **tools and techniques** that account for the different system's components

Recently, **IgProf** was ported for ARM 64 bit and got power profiling capabilities.

ARM keeps showing potential for efficient HTC



Q & A

goncalo.pestana@cern.ch